Ryan Araula

(832) 421-8294 araulaphillipryan@gmail.com https://www.linkedin.com/in/ryanaraula/

PROFESSIONAL SUMMARY

Analytical and detail-driven Computer Science major from the University of Houston, specializing in transforming complex datasets into actionable insights. Experienced in building data pipelines, developing interactive dashboards, and applying machine learning models to real-world problems. Known for strong problem-solving skills, and the ability to quickly learn new tools and technologies.

EDUCATION

University of Houston, Texas
Bachelor of Science in Computer Science *May 2026*

TECHNICAL SKILLS

Programming: SQL, Python, JavaScript
Big Data Frameworks: PySpark, Delta Lake
Cloud Platforms: Azure SQL Server, Azure Data Factory, ADLS Gen2, Azure Synapse, Databricks
Security & Identity: Azure Key Vault, RBAC, Managed Identities
Analytics & BI: Power BI, Jupyter/Colab
Databases: MySQL, PostgreSQL
DevOps: GitHub Actions, Git
Certifications:

- **Microsoft:** Azure AI Fundamentals, Azure Data Fundamentals
- **Databricks:** Academy Accreditation - Databricks Fundamentals
- **DataCamp:** SQL Associate, Data Engineer Associate

PROJECTS / RESEARCH EXPERIENCE - https://rparaula.github.io/

Postal Office Data Warehouse Demo *August 2024 – Present*

- Migrated a legacy MySQL database into an Azure-based data warehouse *(Medallion Architecture)* that provided analytics-ready models for customer demographics, middle-mile package logistics, and root-cause analysis.
- Built batch pipelines using Azure Data Factory to ingest SQL source tables into a bronze ADLS Gen2 container as raw Parquet files, reducing total storage footprint by 76.8%.
- Implemented Databricks/PySpark medallion transformations to convert bronze data into cleaned silver tables and business-ready Delta gold tables, including standardization, deduplication, quality checks, and fact/dimension modeling.
- Built metadata-driven Synapse Serverless SQL data mart views to expose approved fact and dimension tables by reporting domain, improving data governance and management by limiting downstream access to required datasets.
- Integrated Power BI embedded reports into an ASP.NET/C#/Javascript web application that enforced data security through authentication-report access and row-level/object-level security controls.

ETL Pipeline for an Air Quality ML Classifier *January 2025 – Present*

- Created custom data extractors in Python to query third-party air quality/meteorological REST APIs across 97+ Houston ZIP code centroids.
- Automated daily ingestion with GitHub Actions, collecting approximately ~2,368 hourly air quality and meteorological records per day.
- Implemented incremental state tracking and metadata deduplication to ensure idempotent daily runs with automatic crash recovery.

Phishing Email Classifier + Adversarial Machine Learning *August 2025 - December 2025*

- Consolidated multiple email corpora into a clean, standardized multi-class dataset (ham, spam, phishing) for pre- and post-2010 samples.
- Developed and evaluated phishing-detection pipelines using TF-IDF + AdaBoost and DistilBERT + LightGBM (PyTorch/Hugging Face), incorporating RDAP/DNS domain metadata features for comparative robustness analysis.
- Evaluated model performance across datasets and stress-tested pipelines using TextAttack, analyzing model complexity vs. vulnerability.

WORK EXPERIENCE

(Contractor w/ various AI training platforms) *April 2026 – Present*
Data Annotator / Prompt Engineer

- Utilize Reinforcement Learning from Human Feedback (RLHF) to help machine learning models learn more efficiently.
- Ensure models adhere to legal and ethical obligations (i.e., harmlessness, truthfulness, safety, etc.).
- Prompt engineering with various parameters and producing assessments of results.